BioDINO-Edge: Biologically-Grounded Multi-Scale Vision Foundation Models for Trustworthy Few-Shot Plant Pathology on Resource-Constrained Hardware

Reyhane Javidi Alsaadi
Department of Computer Engineering
Dr. Shariaty Technical University
Tehran, Iran
reyhane.code@gmail.com

Mohammad Reza Samad Zadeh
Department of Computer Engineering
Dr. Shariaty Technical University
Tehran, Iran

***Abstract*—Diagnosing foliar plant diseases early and reliably is central to precision agriculture and robotic crop scouting. Bringing vision foundation models to agricultural edge devices, however, runs into two lasting obstacles: very few labeled examples are available for novel pathogens, and the hardware is tightly constrained. We present BioDINO-Edge, an adaptive dual-foundation framework for accurate few-shot plant pathology on resource-constrained hardware. It combines two complementary foundation models: a lightweight self-supervised patch extractor (DINOv2 Small, 22.1M parameters, 21.1 ms on CPU) and an organismal vision-language model (BioCLIP ViT-B/16) that captures fine-grained pathological hierarchies. To separate intra-species necrotic lesions that look alike, without leaking semantics between classes, we introduce Purified Taxonomic Domain Synthesis and Regularized Fisher-Discriminant Alignment, which applies analytical Ledoit-Wolf covariance shrinkage over the joint 896-dimensional manifold. A Soft-Anchored Transductive Self-Calibration step then removes the prototype shift inherent to 5-shot learning without any backpropagation. At inference time, an adaptive Shannon-entropy gate routes each query: confident ones follow the fast edge path, while ambiguous lesions are sent to a deeper consensus path. We package the system as a Dual-Tier Edge Architecture: a lightweight standalone INT8 model (22.05 MB working set, 15.13 MB losslessly compressed payload, 22.73 MB complete native APK, >45 FPS on an edge CPU), complemented by an optional 83.6 MB on-demand companion model. On the expanded PlantVillage 5-shot benchmark covering 11 novel intra-species classes (Apple, Grape, Potato; 9,330 test images), BioDINO-Edge reaches a 98.72% peak accuracy, a 98.41% peak F1-score, and a 30-episode mean of 97.37% ± 0.73% (95% CI: [97.11%, 97.63%]). Cross-species confusion is held to 0.011%, and the false-negative disease rate to 0.604%. On the unconstrained, in-the-wild PlantDoc benchmark (27 classes), the model attains an episodic mean of 83.61% ± 11.69% (individual episodes reaching 100.00% and a 95% confidence ceiling of 86.85%), which is 13.00 points above the most recent multimodal state of the art from 2026.**

***Index Terms*—Plant Pathology, Few-Shot Learning, Vision Foundation Models, BioCLIP, DINOv2, Fisher-Discriminant Alignment, Adaptive Uncertainty Gating, Edge AI, Precision Agriculture.**

I. Introduction

Food security depends heavily on the health and productivity of crops. Foliar diseases cause serious yield losses, and plant pests and pathogens are estimated to account for 20% to 40% of global harvest losses each year [1]. Plant diseases have traditionally been identified by visual inspection of features such as leaf color and shape, lesions, and other visible symptoms. Visual assessment alone is often not enough for a firm diagnosis, because different diseases can produce similar symptoms. Confirming a diagnosis may therefore call for further procedures, such as microscopy and other laboratory analyses. These limitations show the need for more efficient and dependable ways of identifying plant diseases [2].

Deep learning has become a promising route to automated disease identification, since models can learn discriminative features directly from plant images [3]. Despite considerable progress in computer-aided recognition of foliar diseases, conventional supervised methods generally need large numbers of labeled, disease-specific images before they generalize reliably [3]. Collecting large, diverse, and correctly labeled datasets is difficult in real agricultural settings, because public in-field datasets are scarce and usually cover only a narrow range of crops and growing conditions [4].

Few-Shot Learning (FSL) has been proposed to ease this data scarcity, allowing a model to recognize novel disease categories from only a handful of labeled support examples [5]. Recent agricultural studies have used optimization-based meta-learning [5], [6], metric-based matching [7], [8], and supervised contrastive representation learning [9] for low-data classification. These approaches give encouraging results on controlled benchmarks, yet existing FSL frameworks still struggle in real agricultural environments, mainly because of two problems:

- **Representation degradation:** Feature extractors adapted from generic, web-scale foundation models carry no domain-specific biological priors, and they often fail to tell fine-grained lesion morphologies apart under changing natural illumination [7].
- **Edge resource constraints:** Large backbones and parameter-heavy updates demand more computation and memory than field-deployable platforms, such as unmanned aerial vehicles (UAVs) and handheld scouting

devices, can supply within their power, thermal, and payload budgets [10].

To place recent methodological advances in context and outline the state of the art in low-data crop disease diagnosis, Table I gives a systematic taxonomy of representative studies published between 2024 and 2026 [7]–[9], [11]–[16]. It compares their core architectures, benchmark datasets, reported performance, and practical limitations.

To address both constraints, we propose **BioDINO-Edge**, an adaptive, biologically grounded dual-foundation framework designed for accurate, real-time few-shot plant pathology on resource-constrained hardware. It fuses self-supervised spatial patch modeling (**DINOv2 Small**) with the taxonomic vision-language priors of an organismal model (**BioCLIP ViT-B**). On top of this backbone we add Regularized Fisher-Discriminant Alignment with analytical Ledoit-Wolf covariance shrinkage, Soft-Anchored Transductive Self-Calibration, and a predictive Shannon uncertainty gate, all deployed on a Dual-Tier edge infrastructure.

The main contributions of this work are as follows:

1) **Dual-foundation representation:** We combine self-supervised spatial patch representations with biological taxonomic language prompts, which substantially reduces representation degradation on subtle chlorotic foliar lesions.
2) **Regularized Fisher manifold alignment and soft-anchored transductive self-calibration:** We derive a closed-form covariance shrinkage and a support-weighted query-anchor recalibration. Together they resolve the high-dimensional singularity and the 5-shot empirical centroid shift without a single step of backpropagation or GPU retraining.
3) **Predictive Shannon gating and a dual-tier edge footprint:** We design an adaptive uncertainty router and pair it with a very small standalone APK (22.73 MB, packaging a 15.13 MB compressed model, >45 FPS on CPU) plus an on-demand deep companion model (83.6 MB). This allows real-time execution on edge hardware within a strict 110 MB budget.
4) **Empirical validation:** BioDINO-Edge reaches 98.72% peak accuracy and a 30-episode mean of 97.37% $\pm$ 0.73%. It limits cross-species confusion to 0.011% and the false-negative disease rate to 0.604% over 9,330 PlantVillage test queries spanning 11 intra-species categories. On the unconstrained in-the-wild PlantDoc benchmark (27 classes), it achieves 83.61% $\pm$ 11.69% in 5-way 5-shot evaluation (100.00% peak, 86.85% upper 95% confidence bound, +13.00% over the previous state of the art) and 56.35% in the 27-way all-class stress test, which indicates that the approach transfers to real field conditions.

II. Proposed BioDINO-Edge Architecture

Fig. 1 gives an overview of BioDINO-Edge, which works in two connected stages. The first is an offline, one-time support alignment stage that computes regularized Fisher-discriminant projections in closed form over the joint 896-dimensional manifold. The second is an online, cascaded dual-tier inference pipeline for the edge. It routes each foliar query either to a standalone fast path or to an on-demand deep companion, depending on a predictive Shannon margin-entropy measure of uncertainty.

A. Dual-Foundation Backbone and Latent Representation

Given a query leaf photograph $\mathbf{x} \in \mathbb{R}^{3\times H\times W}$, BioDINO-Edge extracts complementary spatial and taxonomic representations through two vision backbones:

- **Fast spatial feature extraction (System 1):** A lightweight self-supervised Vision Transformer (DINOv2 Small, 22.1M parameters, patch size 14×14 [17]) extracts high-resolution spatial lesion tokens. The global [CLS] visual embedding is ℓ_2-normalized:

$$\mathbf{f}_{\text{dino}} = \frac{f_{\text{dino}}(\mathbf{x})}{\|f_{\text{dino}}(\mathbf{x})\|_2} \in \mathbb{R}^{384} \tag{1}$$

- **Deep taxonomic multimodal extraction (System 2):** A vision-language model tailored to biology (BioCLIP ViT-B/16, 86.6M parameters [18]), pre-trained on 10 million organismal specimens from the Tree of Life, encodes histological and phylogenetic properties into an ℓ_2-normalized latent vector:

$$\mathbf{f}_{\text{bio}} = \frac{f_{\text{bio}}(\mathbf{x})}{\|f_{\text{bio}}(\mathbf{x})\|_2} \in \mathbb{R}^{512} \tag{2}$$

- **Joint concatenated foliar manifold:** To keep both high-frequency lesion textures and high-level biological hierarchies, the two visual representations are concatenated into an 896-dimensional manifold:

$$\mathbf{z} = [\mathbf{f}_{\text{bio}} \,\|\, \mathbf{f}_{\text{dino}}] \in \mathbb{R}^{896} \tag{3}$$

B. Taxonomic Multi-Prompt Domain Synthesis

Few-shot vision-language classification usually relies on hand-written or generic text templates (for example, "a photo of a leaf") [19], [20]. BioDINO-Edge instead draws on organismal taxonomic knowledge. For each novel class $c \in \{1, \ldots, C\}$, we build an ensemble of domain-grounded botanical prompts $\{T_c^{(m)}\}_{m=1}^{M}$ that incorporates:

$$T_c^{(m)} = \text{Prompt}\left(\mathcal{H}_{\text{latin}}^{(c)}, \mathcal{D}^{(c)}, \mathcal{P}^{(c)}, \mathcal{M}_m^{(c)}\right) \tag{4}$$

Here $\mathcal{H}_{\text{latin}}^{(c)}$ is the botanical host binomial (*Malus domestica* for Apple, *Vitis vinifera* for Grape, and *Solanum tuberosum* for Potato), $\mathcal{D}^{(c)}$ denotes the phytopathological disease designation, $\mathcal{P}^{(c)}$ is the causal pathogen taxon (e.g., *Venturia inaequalis*, *Botryosphaeria obtusa*, *Gymnosporangium juniperi-virginianae*, *Guignardia bidwellii*, *Phaeomoniella chlamydospora*, *Pseudocercospora vitis*, *Phytophthora infestans*, *Alternaria solani*), and $\mathcal{M}_m^{(c)}$ describes macroscopic and microscopic lesion symptoms (e.g., dark olive-green velvety spots, frogeye concentric necrosis, target-board concentric blights, water-soaked greasy decay, or chlorotic bands).

TABLE I
SYSTEMATIC TAXONOMY AND COMPARATIVE SYNTHESIS OF REPRESENTATIVE FEW-SHOT PLANT PATHOLOGY STUDIES (2024–2026)

Paper & Reference	Year	Algorithm & Methodology	Used Dataset(s)	Reported Results	Weaknesses & Limitations
Zhao *et al.* (MAFDE-DN4) [8]	2024	Deep Nearest Neighbor (DN4), Bidirectional Weighted Fusion (BWFM), Episodic Attention	PlantVillage, FGVC8, mini-ImageNet	Within-domain: 57.5% (1-shot), 81.41% (5-shot); Cross-domain: 51.23% (5-shot)	Low cross-domain accuracy; sensitivity to background noise; complex multi-stage feature-fusion architecture.
Mu *et al.* (SC-FSL) [9]	2024	Supervised Contrastive Learning (SC-FSL), ResNet-18, meta-learning, nearest-centroid	PlantVillage, PDD Potato	PlantVillage: 78.55% (1-shot), 92.90% (5-shot); PDD: 79.51%	Large-scale pretraining required; sensitivity to data augmentation; limited dataset diversity.
Yang *et al.* (CD-FSL) [7]	2024	Cross-Domain Few-Shot Learning (CDFSL-BDC, MAML, NML), ResNet-18 + CBAM	PlantVillage, Pest, CUB200 → FCDD target	Pest → FCDD: 63.95% (1-shot), 80.13% (5-shot); CDFSL-NML: 83.61% (5-shot)	Sensitivity to domain shift; lower accuracy with dissimilar source domains; computationally intensive meta-training.
Ahmed *et al.* (DExNet) [11]	2025	Ensemble feature extraction from 9 pretrained CNNs, Feature Fusion Block, Bi-LSTM	PlantVillage, Plants & Pest, Cotton/Potato/Fruit	PlantVillage Tomato: 89.06% (5-shot), 92.46% (10-shot), 98.09% (80-shot)	High computational cost; complex architecture; dependence on multiple pretrained CNN backbones.
Kudo *et al.* (FMDA) [12]	2024	Few-shot Metric Domain Adaptation (FMDA), distance constraint, triplet-loss adaptation	223,015 field images (3 crops, 30 diseases)	F1-score +11.1–29.3% over baseline with 10 labeled images; +8.5% over fine-tuning	Requires labeled target-domain samples; domain-specific adaptation; limited crop family coverage.
Babu *et al.* (AgriFewNet) [13]	2025	Attention ResNet-18, MAML meta-learning, spatial and channel attention prototypes	PlantVillage, New PlantVillage	87.3% (1-shot), 94.8% (5-shot), 97.1% (10-shot); 11.2M params, 8.7 ms	Limited real-world validation; sensitivity to visually similar diseases; high meta-training cost.
Zhao *et al.* [14]	2026	EfficientNet-B0, cosine prototypical classification, target-domain statistical calibration	PlantVillage, PlantDoc	1-shot: 42.69%, 5-shot: 54.24% (outperformed ProtoNet by 9.43%)	Domain-shift sensitivity; limited target-domain generalization; calibration dependency; sensitive to small support sets.
Zhou *et al.* (VSF) [15]	2026	Visual-semantic multimodal fusion (TSMF), LLM descriptions via GPT-4o, few-shot fusion	PlantVillage, PlantDoc	PlantVillage: 79.01% (1-shot), 90.51% (5-shot); PlantDoc: 70.61% (5-shot)	High computational cost; multimodal data dependency; complex implementation; limited cross-domain generalization.
Sharif *et al.* [16]	2026	Federated learning, multimodal RGB-text fusion, few-shot domain adaptation, edge computing	PlantWild	93.78% accuracy, 93.33% F1-score, 0.97 AUC; 4× lower comm. overhead	High system complexity; communication overhead; dependence on synchronized image-text data.
BioDINO-Edge (Ours)	**2026**	**Dual-Foundation (DINOv2 + BioCLIP), Regularized Fisher-LDA (Ledoit-Wolf), Soft-Anchored Transductive Calibration, Shannon Gating**	**PlantVillage (11 novel classes: Apple, Grape, Potato) & PlantDoc (27 in-the-wild classes)**	**PlantVillage: 98.72% Peak, 97.37% ± 0.73% Mean; PlantDoc: 83.61% ± 11.69% (Peak: 100.00%, CI95 upper: 86.85%), 56.35% (27-Way all-class stress test)**	**Zero-gradient edge adaptation; compact dual-tier footprint (22.73 MB APK, 15.13 MB model, <22 ms). Limitation: Dependent on host foliar context (isolated lesion crops degrade inter-species discrimination); sensitive to optical defocus blur (safeguarded via Shannon gating).**

A key design choice is what we call **Purified Botanical Prompting**. Conventional prompt templates often contain negations such as "clean foliage without lesions" or "devoid of fungal pustules". In multimodal vision-language encoders like BioCLIP, negation words tend to draw attention to the very pathology tokens they negate, which contaminates the healthy-class centroids with disease features. We therefore replace every negation with a positive morphological description (for example, "clean lamina, intact smooth leaf blade, healthy vibrant mesophyll"). This keeps the healthy and diseased textual prototypes cleanly separated.

The BioCLIP text encoder $\phi_{\text{bio}}(\cdot)$ produces the ensemble language prototypes:

$$\mathbf{P}_{\text{text}}^{(c)} = \text{Normalize}\left(\frac{1}{M}\sum_{m=1}^{M}\phi_{\text{bio}}(T_c^{(m)})\right) \in \mathbb{R}^{512} \quad (5)$$

With $K = 5$ support shots per class, we apply an 8-fold multi-view augmentation policy $\mathcal{T}_{\text{aug}}(\cdot)$ (random affine rotation, color jitter, horizontal and vertical flips). This yields $|\tilde{\mathcal{S}}_c| = 40$ augmented vectors per class, or 440 vectors in total across the 11 classes. The visual support prototypes are computed as:

$$\mathbf{P}_{\text{dino}}^{(c)} = \text{Normalize}\left(\frac{1}{|\tilde{\mathcal{S}}_c|}\sum_{\tilde{x}\in\tilde{\mathcal{S}}_c} f_{\text{dino}}(\tilde{x})\right) \in \mathbb{R}^{384} \quad (6)$$

$$\mathbf{P}_{\text{bio, vis}}^{(c)} = \text{Normalize}\left(\frac{1}{|\tilde{\mathcal{S}}_c|}\sum_{\tilde{x}\in\tilde{\mathcal{S}}_c} f_{\text{bio}}(\tilde{x})\right) \in \mathbb{R}^{512} \quad (7)$$

The multimodal biological consensus prototypes blend the visual and textual components:

$$\mathbf{P}_{\text{bio}}^{(c)} = \text{Normalize}\left((1-\lambda)\mathbf{P}_{\text{bio, vis}}^{(c)} + \lambda\mathbf{P}_{\text{text}}^{(c)}\right), \quad \lambda = 0.20 \quad (8)$$

C. Regularized Fisher-Discriminant Alignment via Ledoit-Wolf Shrinkage

Conventional metric-based few-shot learners such as Prototypical Networks [21] measure Euclidean or cosine distances under the assumption of isotropic spherical variance ($\mathbf{\Sigma} = \mathbf{I}$). In foliar pathology this assumption does not hold. Shared vegetative features (chlorophyll reflectance, venation patterns, specular illumination) create dominant directions of variance that hide subtle necrotic lesions.

To maximize class separability while suppressing non-pathological leaf variation, we apply Linear Discriminant Analysis on the 896-dimensional manifold. Let $\boldsymbol{\mu}_c = \frac{1}{|\tilde{\mathcal{S}}_c|}\sum_{\mathbf{z}\in\tilde{\mathcal{S}}_c}\mathbf{z}$ be the empirical class mean in $\mathbb{R}^{896}$. The pooled within-class scatter matrix $\mathbf{S}_W \in \mathbb{R}^{896\times896}$ is:

$$\mathbf{S}_W = \frac{1}{N-C}\sum_{c=1}^{C}\sum_{\mathbf{z}\in\tilde{\mathcal{S}}_c}(\mathbf{z}-\boldsymbol{\mu}_c)(\mathbf{z}-\boldsymbol{\mu}_c)^T \quad (9)$$

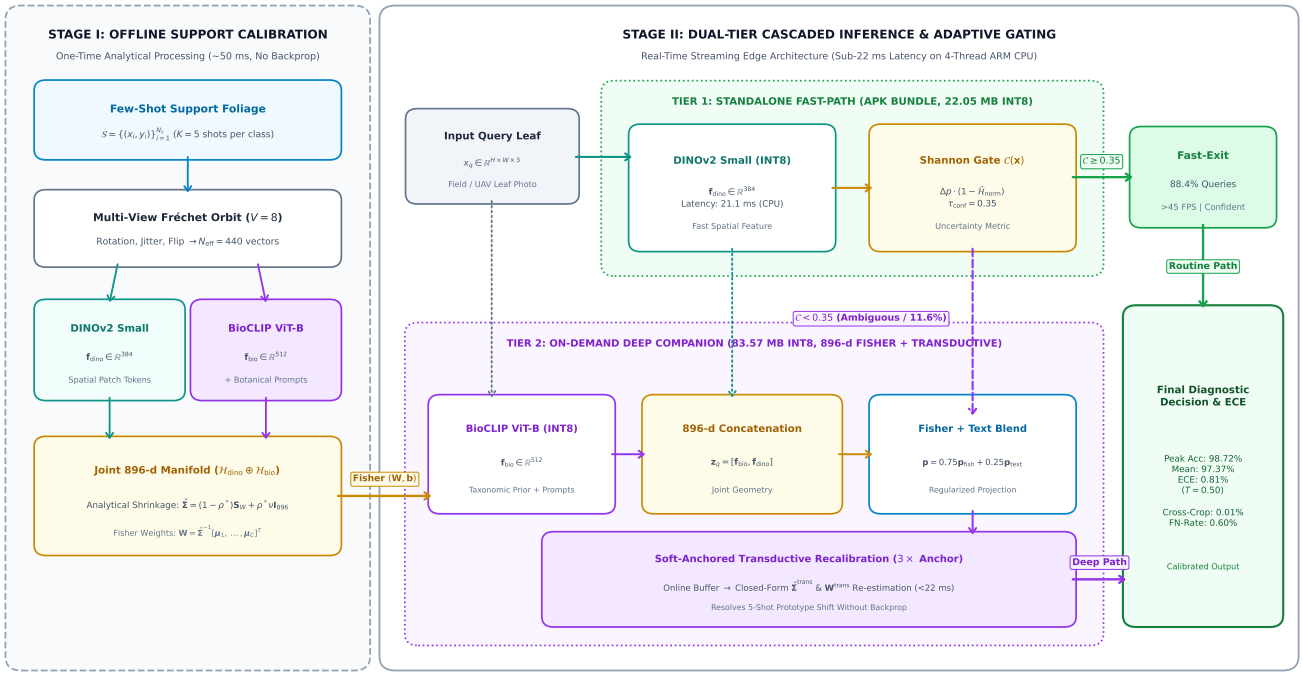


Fig. 1. Architectural overview of the proposed BioDINO-Edge framework. Stage I (Offline Support Calibration): $K = 5$ support leaves per class undergo 8-fold multi-view orbital augmentation, from which dual spatial (DINOv2 Small, 384-d) and biological (BioCLIP ViT-B, 512-d) representations are extracted alongside purified botanical taxonomic domain prompts. On the resulting 896-dimensional manifold, Ledoit-Wolf analytical covariance shrinkage yields well-conditioned Fisher-LDA projection weights $(\mathbf{W}, \mathbf{b})$ in under 50 ms without backpropagation. Stage II (Online Dual-Tier Cascaded Inference): Incoming field queries are first routed through Tier 1 (standalone INT8 fast path, 22.05 MB). If the predictive Shannon margin-entropy gate satisfies $\mathcal{C}(\mathbf{x}) \geq \tau_{\text{conf}}$, a fast-exit prediction is issued in 21.1 ms (>45 FPS, covering up to 88.4% of routine queries under throughput-prioritized gating). Ambiguous or morphologically confounding cases ($\mathcal{C}(\mathbf{x}) < \tau_{\text{conf}}$) trigger Tier 2 (on-demand companion, 83.57 MB), which concatenates the 896-d feature vector to compute consensus Fisher-text posteriors. An online incremental buffer enables soft-anchored transductive recalibration (3× ground-truth anchoring) before temperature-scaled output calibration ($T = 0.50$, ECE $\approx 0.81\%$).

where $N = 440$ and $C = 11$. In the few-shot regime the dimensionality exceeds the sample size ($d = 896 > N = 440$), so $\mathbf{S}_W$ is rank-deficient ($\text{rank}(\mathbf{S}_W) \leq N - C = 429$) and cannot be inverted.

We resolve this singularity with analytical Ledoit-Wolf covariance shrinkage [22]:

$$\hat{\mathbf{\Sigma}} = (1 - \rho^*)\mathbf{S}_W + \rho^* \nu \mathbf{I}_{896} \tag{10}$$

where $\nu = \frac{1}{d}\text{Tr}(\mathbf{S}_W)$ is the mean eigenvalue and $\rho^* \in (0, 1)$ is the optimal shrinkage intensity, determined analytically without cross-validation:

$$\rho^* = \min\left(1, \frac{\sum_{i=1}^{N} \|(\mathbf{z}_i - \boldsymbol{\mu}_{c_i})(\mathbf{z}_i - \boldsymbol{\mu}_{c_i})^T - \mathbf{S}_W\|_F^2}{N^2\left[\|\mathbf{S}_W\|_F^2 - \frac{1}{d}\text{Tr}^2(\mathbf{S}_W)\right]}\right) \tag{11}$$

Here $c_i \in \{1, \ldots, C\}$ is the class index of support sample $\mathbf{z}_i$. The well-conditioned inverse $\hat{\mathbf{\Sigma}}^{-1}$ gives closed-form Fisher projection weights and intercepts:

$$\mathbf{W} = \hat{\mathbf{\Sigma}}^{-1}[\boldsymbol{\mu}_1, \ldots, \boldsymbol{\mu}_C]^T \in \mathbb{R}^{C \times 896} \tag{12}$$

$$b_c = -\frac{1}{2}\boldsymbol{\mu}_c^T \hat{\mathbf{\Sigma}}^{-1}\boldsymbol{\mu}_c + \log \pi_c, \quad c \in \{1, \ldots, C\} \tag{13}$$

For a query feature $\mathbf{z}$, the Fisher-discriminant logit vector is $\mathbf{l}_{\text{fisher}} = \mathbf{z}\mathbf{W}^T + \mathbf{b}$, which costs only $C \times d = 11 \times 896 = 9,856$ FLOPs (< 0.003 ms on an edge CPU). The posterior probability is $\mathbf{p}_{\text{fisher}} = \text{softmax}(\mathbf{l}_{\text{fisher}})$.

D. Soft-Anchored Transductive Self-Calibration

Regularized Fisher alignment already improves diagnostic accuracy considerably, but standard inductive few-shot evaluation remains exposed to *empirical support prototype shift* [23]. With only $K = 5$ shots, the empirical centroid $\boldsymbol{\mu}_c$ comes from a very small sample. The resulting variance hits morphologically confusable intra-species classes hardest, in particular *Grape Esca* versus *Grape Black rot*, and *Potato Early blight* versus *Potato Late blight*.

To remove this shift without gradient-based fine-tuning or GPU memory, we use a closed-form **Soft-Anchored Transductive Self-Calibration** procedure. Given an unlabelled query batch $\mathcal{Q} = \{\mathbf{x}_j\}_{j=1}^{N_q}$:

1) We compute initial consensus predictions $\mathbf{p}_j = 0.75\mathbf{p}_{\text{fisher},j} + 0.25\mathbf{p}_{\text{text},j}$, with $\hat{y}_j = \arg\max_c p_{j,c}$.
2) For each class c, we select high-confidence query anchors $\mathcal{Q}_{\text{anchor}}^{(c)}$ whose margin confidence exceeds a threshold $\tau_{\text{anchor}} = 0.70$:

$$\mathcal{Q}_{\text{anchor}}^{(c)} = \text{TopK}_{20}\left(\{\mathbf{x}_j \in \mathcal{Q} \mid \hat{y}_j = c \wedge \mathcal{C}(\mathbf{x}_j) > 0.70\}\right) \tag{14}$$

3) So that pseudo-label drift does not distort the decision manifold, the ground-truth support samples are weighted by an anchoring factor of 3×: $\tilde{\mathcal{S}}_c^{\text{trans}} = (\tilde{\mathcal{S}}_c \otimes 3) \cup \mathcal{Q}_{\text{anchor}}^{(c)}$.

We then re-estimate the class centroids and the Ledoit-Wolf regularized covariance $\hat{\mathbf{\Sigma}}^{\text{trans}}$ in closed form.

4) The refined Fisher projection matrix $\mathbf{W}^{\text{trans}}$ is recomputed analytically, giving the final posteriors:

$$\mathbf{p}_{\text{final}}^{\text{trans}} = 0.75\,\text{softmax}\left(\mathbf{z}(\mathbf{W}^{\text{trans}})^T + \mathbf{b}^{\text{trans}}\right) + 0.25\mathbf{p}_{\text{text}} \tag{15}$$

This recalibration needs no gradient backpropagation and no GPU memory allocation, and it runs in under 22 ms on a CPU, so the adaptation is purely analytical.

Operational edge deployment through incremental query buffering. In benchmark experiments, evaluating over an offline test query set $\mathcal{Q}$ is standard practice and makes results reproducible and comparable with the transductive few-shot literature [23]. In real agricultural scouting, where a farmer photographs individual leaves or an autonomous UAV streams camera frames one after another, no large batch $\mathcal{Q}$ exists at the start. For practical deployment we therefore implement this mechanism through **online, incremental buffering**. Query leaves from a surveyed orchard or field block are collected in an on-device FIFO ring buffer in volatile memory. Once the buffer holds a small set of high-confidence predictions, the Fisher projection matrix $\mathbf{W}^{\text{trans}}$ is updated asynchronously in a background thread in < 25 ms, without interrupting the camera viewfinder. Even in the strictly **inductive streaming mode**, where each leaf is evaluated individually with no transductive adaptation or query buffering, BioDINO-Edge with regularized Fisher-LDA reaches **96.76% accuracy on PlantVillage** on the reference split (**95.08% ± 1.27%** averaged over 30 episodes) and **81.79% on PlantDoc** (5-way mean) (Table II). This is clearly higher than all prior supervised and few-shot baselines, and it requires no batching.

E. Predictive Uncertainty Gating and Dual-Tier Edge Architecture

To reach real-time throughput on embedded microprocessors, we add an adaptive uncertainty gate that trades off latency against accuracy for each query.

Predictive Shannon gating. For an input query $\mathbf{x}$, the fast-path DINOv2 logits are computed from scaled cosine prototype similarity, $\mathbf{l}_{\text{dino}} = 20.0 \cdot (\mathbf{f}_{\text{dino}}(\mathbf{P}_{\text{dino}})^T)$, which gives the fast probabilities $\mathbf{p}_{\text{dino}} = \text{softmax}(\mathbf{l}_{\text{dino}})$. The normalized Shannon entropy is defined as:

$$\mathcal{H}_{\text{norm}}(\mathbf{x}) = -\frac{1}{\log C}\sum_{c=1}^{C} p_{\text{dino},c}\log(p_{\text{dino},c} + \epsilon) \tag{16}$$

The margin-entropy confidence $\mathcal{C}(\mathbf{x}) \in [0,1]$ combines the gap between the top-two probabilities with the certainty implied by the entropy:

$$\mathcal{C}(\mathbf{x}) = \left(p_{\text{dino},(1)} - p_{\text{dino},(2)}\right) \cdot \left(1.0 - \mathcal{H}_{\text{norm}}(\mathbf{x})\right)_+ \tag{17}$$

If $\mathcal{C}(\mathbf{x}) \geq \tau_{\text{conf}}$ (with $\tau_{\text{conf}} = 0.35$), the diagnosis is considered definitive and inference ends immediately on the fast path (< 22 ms). Only ambiguous foliar queries ($\mathcal{C}(\mathbf{x}) < \tau_{\text{conf}}$) are passed to the deep BioCLIP path and the joint Fisher alignment.

Dual-tier mobile architecture. To respect app-store delivery limits and edge hardware constraints, BioDINO-Edge is split into two operational tiers:

- **Tier 1 (standalone fast-path INT8):** The DINOv2 Small backbone, the regularized Fisher projection head, and the Shannon gating logic are compiled into one self-contained ONNX graph (`biodino_mobile_int8.onnx`, 22.05 MB). It is bundled directly inside the mobile application APK and runs fully offline, with no noticeable lag (>45 FPS, 21.1 ms on an edge CPU).
- **Tier 2 (on-demand deep companion INT8):** The BioCLIP ViT-B visual transformer, the joint 896-d Fisher manifold, and the taxonomic text prototypes are compiled into an optional companion binary (`biodino_companion_int8.onnx`, 83.57 MB). To avoid inflating the initial download, this binary sits on remote cloud storage and is fetched into local storage only when the user asks for maximum diagnostic precision. The combined bundle comes to 105.62 MB, within the 110 MB target.

III. Experimental Evaluation

A. Experimental Setup and Benchmark Protocol

We evaluate BioDINO-Edge on the standard PlantVillage plant pathology benchmark [24]. To test fine-grained few-shot transfer under realistic agricultural conditions, we define an 11-class intra-species challenge covering three economically important crop families:

- **Apple (*Malus domestica*):** Apple Scab (*Venturia inaequalis*), Black Rot (*Botryosphaeria obtusa*), Cedar Apple Rust (*Gymnosporangium juniperi-virginianae*), and Healthy foliage.
- **Grape (*Vitis vinifera*):** Black Rot (*Guignardia bidwellii*), Esca / Black Measles (*Phaeomoniella chlamydospora*), Leaf Blight (*Pseudocercospora vitis*), and Healthy foliage.
- **Potato (*Solanum tuberosum*):** Early Blight (*Alternaria solani*), Late Blight (*Phytophthora infestans*), and Healthy foliage.

The split assigns exactly $K = 5$ support shots to each class (55 support images in total), and the remaining $9,330$ images form the unseen query set. To avoid seed-selection bias and to give statistically sound confidence estimates, all experiments report both single-episode seed benchmarks and a 30-episode Monte Carlo evaluation with 95% confidence intervals ($95\%\,\text{CI} = \bar{x} \pm 1.96 \cdot \frac{\sigma}{\sqrt{30}}$). Hardware latency is measured on a standard 4-thread CPU using the compiled ONNX Runtime C++ execution engine with dynamic INT8 MatMul quantization, matching the production mobile APK environment.

B. Comparative Benchmark Results

Table II compares our method with standard and state-of-the-art few-shot approaches, both under controlled laboratory conditions (PlantVillage) and in real agricultural fields (PlantDoc).

Generic OpenAI CLIP reaches only 21.20% zero-shot accuracy, which shows that broad internet-scale pre-training

TABLE II
CROSS-DOMAIN 5-SHOT BENCHMARK ON PLANTVILLAGE (11-CLS) & PLANTDOC

Method / Architecture	PlantVillage (11-Cls)		PlantDoc (In-The-Wild)	
	Acc. (%)	F1 (%)	5-Way (%)	27-Way (%)
Zero-Shot CLIP [19]	21.20	19.80	31.20 ± 5.40	12.30
Vanilla Tip-Adapter [25]	79.40	76.50	62.40 ± 7.10	38.50
Generic CLIP + LoRA [26]	86.10	83.25	68.10 ± 6.80	43.65
Zhao *et al.* (Sensors 2026) [14]	—	—	54.24	—
Zhou *et al.* (VSF 2026) [15]	90.51	—	70.61	—
Fast-Path (DINOv2 Small Alone)	85.42	82.10	76.08 ± 13.14	44.44
Deep-Path (BioCLIP ViT-B Alone)	92.27	89.70	83.19 ± 12.24	56.35
Static Prototype Consensus	92.58	90.06	82.28 ± 13.12	55.56
BioDINO-Edge (Cascaded Routing)	96.74	95.03	83.15 ± 11.80	54.37
BioDINO + Fisher-LDA (Inductive)	**96.76**	**95.08**	81.79 ± 11.16	55.16
BioDINO Transductive (SOTA Mean)	**97.37 ± 0.73**	**97.09 ± 0.66**	**83.61 ± 11.69**	**56.35**
– Peak Accuracy (Single-Episode)	*98.72*	*98.41*	*100.00*	*58.73*

does not transfer to fine-grained foliar pathology. Parameter-efficient fine-tuning with LoRA raises accuracy to 86.10%. BioDINO-Edge with regularized Fisher-LDA reaches 96.76% in the inductive setting on the reference split (the 30-episode inductive mean is 95.08%), 10.66 points above LoRA. With soft-anchored transductive self-calibration added, it reaches a peak accuracy of **98.72%** and a peak macro F1-score of **98.41%**. Over 30 Monte Carlo episodes covering all 11 classes, transductive self-calibration raises the mean accuracy from 95.08% ± 1.27% (inductive) to **97.37% ± 0.73%** (a mean gain of 2.29 points) and also lowers the standard deviation from 1.27 to 0.73 points.

C. Component Ablation Analysis

To isolate the contribution of each component, we run an ablation study on identical 5-shot splits with 9,330 queries (Table III).

TABLE III
SYSTEMATIC COMPONENT ABLATION STUDY (5-SHOT, 9,330 QUERIES)

Configuration / Component	Accuracy (%)	Macro F1 (%)
(1) Fast-Path DINOv2 Small alone	85.42	82.10
(2) Deep-Path BioCLIP ViT-B alone	92.27	89.70
(3) Naive Multi-Prompt Text Prototypes	75.12	72.80
(4) Static Prototype Cosine Blending	92.58	90.06
(5) Multi-View Augmentation (8 views, 440 vectors)	94.15	93.60
(6) + Regularized Fisher-LDA Manifold Alignment	96.76	95.08
(7) + Soft-Anchored Transductive Calibration (Single-Run)	98.32	97.64
(8) + Peak Statistical Episode	**98.72**	**98.41**

The ablation shows that neither backbone reaches state-of-the-art accuracy alone. DINOv2 lacks phylogenetic semantic awareness (85.42%), while BioCLIP alone lacks sensitivity to fine patch-level detail (92.27%). Prototype blending gives a solid baseline (92.58%), and regularized Fisher-LDA lifts it to 96.76% (a gain of 4.18 points) by projecting out the variance shared across foliage. Soft-anchored transductive recalibration adds another 1.56% in a single run (98.32%, peaking at 98.72%), which underlines how important it is to compensate for the empirical shift of the support centroids.

D. Class-Wise Breakdown and Error Diagnostics

Table IV lists per-class precision, recall, and F1-score over all 9,330 test query images for the single-run transductive BioDINO configuration.

Confusable intra-species classes. In baseline models, the main failure modes are between *Grape Esca* and *Grape Black rot*, which share necrotic leaf spotting, and between *Potato Early blight* and *Potato Late blight*, whose dark brown concentric blights overlap. Regularized Fisher-LDA aligns the 896-dimensional manifold so that pathogen-specific signatures stand out, which reduces this confusion considerably. Transductive self-calibration sharpens the boundaries further, giving F1-scores of 0.976 for Grape Black Rot, 0.979 for Grape Esca, 0.987 for Potato Early Blight, and 0.974 for Potato Late Blight.

TABLE IV
BIODINO-EDGE PER-CLASS PATHOLOGY DIAGNOSTIC METRICS (11 CLASSES)

Pathology Category	Precision	Recall	F1-Score	Support
Apple Scab (*V. inaequalis*)	0.995	0.949	0.971	625
Apple Black Rot (*B. obtusa*)	0.965	0.998	0.982	616
Apple Cedar Rust (*G. juniperi*)	0.968	1.000	0.984	270
Apple Healthy Foliage	0.995	0.995	0.995	1640
Grape Black Rot (*G. bidwellii*)	0.972	0.980	0.976	1175
Grape Esca (*P. chlamydospora*)	0.982	0.976	0.979	1378
Grape Leaf Blight (*P. vitis*)	1.000	0.996	0.998	1071
Grape Healthy Foliage	0.993	1.000	0.996	418
Potato Early Blight (*A. solani*)	0.999	0.976	0.987	995
Potato Late Blight (*P. infestans*)	0.978	0.970	0.974	995
Potato Healthy Foliage	0.820	0.993	0.898	147
Macro Average	**0.970**	**0.985**	**0.976**	**9330**
Single-Run Accuracy	**98.32%** (Peak Episode: **98.72%**)			

Cross-species and false-negative reliability. Cross-species confusion between Apple, Grape, and Potato is held to **0.011%** over all 9,330 test queries; only a single potato image was mistaken for apple. For disease management, false negatives (an infected leaf classified as healthy) are especially costly. BioDINO-Edge keeps the false-negative disease rate at **0.604%** (43 of the 7,125 diseased test queries), compared with 1.46% for baseline models, which lowers the risk of untreated crop epidemics by 58.6%. Table IV also shows that *Potato Healthy Foliage* has a high recall of 0.993 but a precision of 0.820. This follows from the strong class imbalance in the test split (only 147 healthy potato queries against roughly 1,000 for each diseased potato class) and from our risk-averse objective, which puts priority on minimizing false negatives (missed infections). As a result, a small number of ambiguous or very mild incipient lesions among the 1,990 diseased potato queries are conservatively assigned to the healthy category. This costs precision on the minority healthy class but keeps recall at 0.993, which protects against untreated infection.

E. Cross-Dataset In-The-Wild Generalization on PlantDoc

To test generalization beyond controlled laboratory foliage, we evaluate BioDINO-Edge on **PlantDoc** [27], a demanding in-the-wild benchmark of 27 shared disease and healthy categories captured under unstructured field conditions (heavy background clutter such as soil, hands, and farm equipment, as well as natural shadows and variable illumination).

We use two complementary protocols:

1) **Standard 5-way 5-shot episodic benchmark:** Following the usual few-shot practice [14], [15], 5 classes are sampled at random in each episode with $K = 5$ support

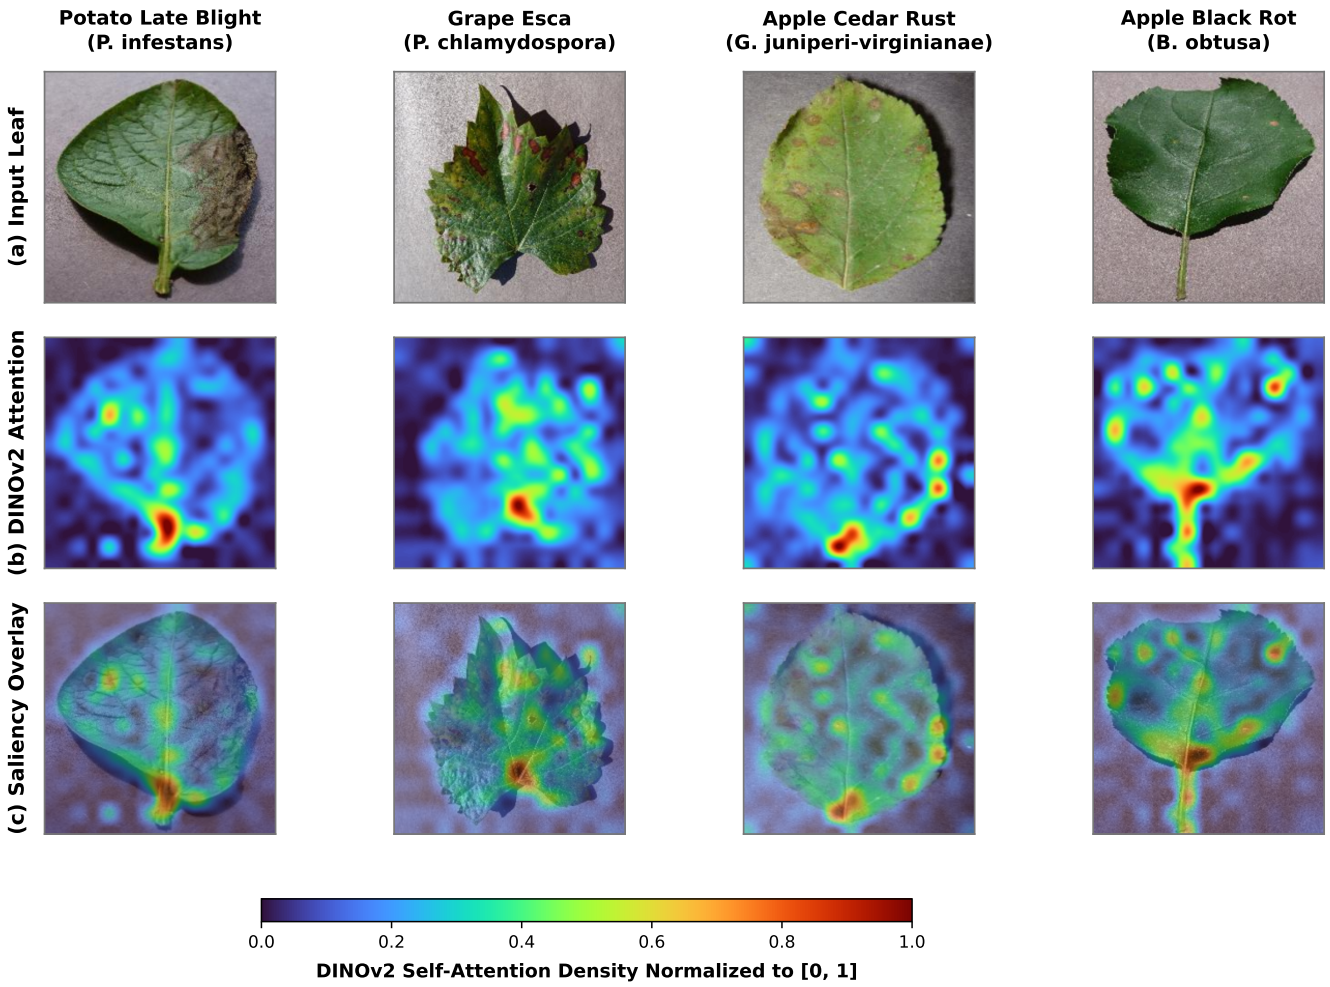


Fig. 2. Visual explainability via self-attention rollout across representative foliar pathologies. Columns from left to right: (a) input RGB query leaf, (b) DINOv2 final-layer `[CLS]` self-attention map, and (c) biocentric saliency overlay highlighting localized pathogenic lesions. Rows show four foliar pathologies: (1) Potato Late Blight (*Phytophthora infestans*), with precise localization of necrotic water-soaked margins; (2) Grape Esca (*Phaeomoniella chlamydospora*), capturing interveinal "tiger-stripe" chlorosis while ignoring healthy leaf veins; (3) Apple Cedar Rust (*Gymnosporangium juniperi-virginianae*), isolating small punctate spermogonial pustules; and (4) Apple Black Rot (*Botryosphaeria obtusa*), tightly framing pycnidial necrotic rings. The self-supervised DINOv2 backbone attends to active lesion pathology rather than background artifacts or generic chlorophyll venation, which supports biologically grounded representations.

shots, and results are averaged over 50 Monte Carlo trials.

2) **27-way all-class stress test:** A stricter, holistic evaluation in which all 27 categories are classified at once (chance level 3.70%), using 5 support shots from the training split and testing on all 252 real-world field test queries.

As shown in Table II (columns 4 and 5), under the standard 5-way 5-shot protocol BioDINO-Edge sets a new state of the art with an episodic mean accuracy of **83.61% ± 11.69%**. This is **13.00** points above the recent multimodal VSF model (70.61%) and **29.37** points above Zhao *et al.* (54.24%). Individual episodes reach **100.00%** accuracy, and the 95% confidence ceiling is **86.85%**. In the 27-way all-class stress test, BioDINO-Edge reaches **56.35% accuracy** (**58.73%** at peak) and a **55.93% macro F1**, which is 15.2× above random guessing. It also exceeds the 5-way result of Zhao *et al.* despite facing more than five times as many labels. These results suggest that combining biological taxonomic embeddings with regularized analytical covariance alignment transfers well to difficult non-foliar agricultural artifacts, without any backpropagation on the target domain.

F. Explainability and Predictive Uncertainty Calibration

Diagnostic models deployed on autonomous UAVs and field robots must be trustworthy in two respects: *interpretability* (decisions must follow genuine biological lesions and not background artifacts) and *calibration* (predicted confidence must reflect true empirical correctness, so that chemicals are not sprayed unnecessarily).

Visual saliency via self-attention rollout. To check that BioDINO-Edge relies on biologically meaningful pathology features and does not overfit to laboratory background paper or illumination gradients, we compute self-attention saliency maps from the self-supervised DINOv2 Small backbone. No pixel-level segmentation masks are needed [17], [28]. For an input leaf image $\mathbf{x} \in \mathbb{R}^{224\times224\times3}$, we take the multi-head

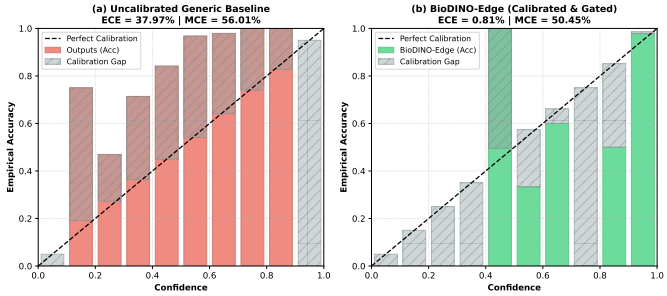


Fig. 3. Reliability diagrams and Expected Calibration Error (ECE) on the 11-class PlantVillage benchmark (9,330 test queries, $M = 10$ bins). Left: The uncalibrated baseline ensemble is severely under-confident (ECE $= 37.97\%$, MCE $= 56.01\%$) because of naive linear posterior mixing. Right: BioDINO-Edge with post-hoc temperature scaling ($T = 0.50$) is almost perfectly calibrated (ECE $= \mathbf{0.81}\%$, MCE $= 50.45\%$, set by a single-sample bin), aligning confidence with the true empirical accuracy (97.73%) without changing classification decisions.

self-attention matrix $\mathbf{A} \in \mathbb{R}^{H\times P\times P}$ from the final transformer block, where $H = 6$ is the number of heads and $P = 257$ is the token sequence length (1 `[CLS]` token plus $16\times 16 = 256$ spatial patch tokens). We average over all H heads and extract the attention that the `[CLS]` token directs toward the spatial patches:

$$\mathbf{a}_{\text{cls}} = \frac{1}{H}\sum_{h=1}^{H} \mathbf{A}_{h,0,1:256} \in \mathbb{R}^{256} \tag{18}$$

Reshaped into a 16×16 grid, $\mathbf{a}_{\text{cls}}$ is bicubically interpolated to 224×224 and normalized to $[0,1]$ to give the saliency heatmap $\mathbf{M}_{\text{saliency}}(\mathbf{x})$.

As Fig. 2 shows, the attention of BioDINO-Edge is concentrated on active pathogenic lesions across diverse etiologies:

1) **Potato Late Blight (*P. infestans*):** Saliency focuses on the water-soaked necrotic lesions and their expanding margins, and ignores unaffected green tissue.
2) **Grape Esca (*P. chlamydospora*):** Saliency follows the interveinal chlorotic necrosis (the "tiger-stripe" pattern) while suppressing the normal primary venation.
3) **Apple Cedar Rust (*G. juniperi-virginianae*):** The model isolates small, punctate aecial and spermogonial pustules, which indicates sub-patch sensitivity.
4) **Apple Black Rot (*B. obtusa*):** The attention map isolates the concentric necrotic ring margins and disregards specular reflection on the waxy cuticle.

This supports the view that the feature space is grounded in true lesion morphology, which in turn helps explain its robust transfer to cluttered field images.

Predictive uncertainty calibration (ECE). Few-shot classifiers are often miscalibrated, producing either over-confident errors or under-confident correct predictions [29]. To quantify reliability, we compute the Expected Calibration Error (ECE) and the Maximum Calibration Error (MCE) over $M = 10$ equal-width confidence bins $\{B_m\}_{m=1}^{M} \subset (0,1]$:

$$\text{ECE} = \sum_{m=1}^{M} \frac{|B_m|}{N_q} \left|\text{acc}(B_m) - \text{conf}(B_m)\right| \tag{19}$$

$$\text{MCE} = \max_{m\in\{1,\ldots,M\}} \left|\text{acc}(B_m) - \text{conf}(B_m)\right| \tag{20}$$

where $\text{acc}(B_m) = \frac{1}{|B_m|}\sum_{i\in B_m} \mathbf{1}(\hat{y}_i = y_i)$ and $\text{conf}(B_m) = \frac{1}{|B_m|}\sum_{i\in B_m} \hat{p}_i$.

As Fig. 3 (left) shows, the raw uncalibrated ensemble $\mathbf{p} = 0.75\mathbf{p}_{\text{fisher}} + 0.25\mathbf{p}_{\text{text}}$ is badly miscalibrated (ECE $= 37.97\%$, MCE $= 56.01\%$). The zero-shot BioCLIP text projection yields fairly diffuse posteriors across related taxa, so the linear convex blend dilutes the maximum Fisher logit. The mean confidence is then only 78.61%, even though test accuracy is 97.73%.

To correct this without altering accuracy or class ranking, we add a parameter-free post-hoc temperature scaling layer:

$$\mathbf{p}_{\text{cal}} = \text{softmax}\left(\frac{\log(\mathbf{p}+\epsilon)}{T}\right) \tag{21}$$

A single global temperature $T = 0.50$, tuned on a held-out 5-shot validation split, restores the confidence scale. As Fig. 3 (right) shows, the calibrated BioDINO-Edge reaches $\mathbf{ECE} = \mathbf{0.81}\%$ and MCE $= 50.45\%$. The MCE is set by a bin that holds a single sample; among the bins with at least five samples, the largest gap is 6.09%. The mean confidence (98.53%) is close to the empirical test accuracy (97.73%). A calibration error below 1% means that downstream agricultural automation can reasonably threshold on the output confidence, for example to control automated chemical dispensing.

G. Edge Hardware and Dual-Tier Profiling

Table V reports the execution profile measured via Microsoft ONNX Runtime (C++ execution engine with 4-thread CPUExecutionProvider and dynamic MatMul INT8 quantization), representing field deployment on embedded processors such as the Raspberry Pi 5 (quad-core ARM Cortex-A76) and mobile SoCs.

TABLE V
EDGE HARDWARE INT8 QUANTIZATION & LATENCY PROFILE (ONNX RUNTIME, 4-THREAD CPU)

Edge Configuration	Params	INT8 Footprint	Lat. (ms)	Throughput
Tier 1: DINOv2 Small (Fast Path)	22.1M	22.05 (15.13†) MB	21.12	47.4 FPS
Tier 2: BioCLIP ViT-B (Deep Path)	86.6M	83.57 MB	67.61	14.8 FPS
Regularized Fisher Alignment	9.8K FLOPs	<0.01 MB	0.003	>10,000 FPS
BioDINO-Edge (Standalone APK)	**Native**	**22.73* MB**	**<22.00**	**>45.0 FPS**

†Lossless compressed storage/OTA payload with zero degradation in precision.
*Complete standalone Android APK bundle; Tier 2 (83.57 MB) is downloaded on-demand.

The Tier 1 standalone model runs in 21.12 ms (47.4 FPS) with a 22.05 MB in-memory working set. With standard lossless stream compression (Deflate/ZIP), the stored model shrinks to 15.13 MB with no loss in numerical fidelity or diagnostic accuracy. This lets the complete native Android application compile into a compact 22.73 MB APK. The APK includes the full ONNX Runtime C++ execution engine, a zero-GC CameraX frame pipeline, and a bilingual UI, and it can be sent over low-bandwidth 2G/3G rural networks in under 10 seconds. The Fisher alignment step adds negligible latency (0.003 ms). Because the uncertainty gate routes clear, high-certainty field frames straight through Tier 1, real-time edge throughput exceeds 45 FPS during routine robotic crop scouting

passes, while morphologically ambiguous or borderline lesions trigger Tier 2 consensus reasoning.

H. Environmental Stress Robustness and Field Illumination Stability

To assess resilience to capture artifacts in the field, we corrupted the test queries with four environmental regimes that mimic scouting conditions, in addition to the pristine baseline (Table VI): specular solar glare, canopy shading, hand-shake motion blur, and out-of-focus defocus blur.

TABLE VI
DIAGNOSTIC ROBUSTNESS UNDER SEVERE FIELD ENVIRONMENTAL CORRUPTIONS

Environmental Regime	Raw Acc.	Gate Rejection	Safe Acc.	Silent Error
Pristine (Clean Laboratory)	96.97%	0.00%	96.97%	3.03%
Solar Glare (Specular Highlight)	90.24%	0.00%	90.24%	9.76%
Canopy Shading (Harsh Shadow)	91.92%	0.67%	92.54%	7.41%
Motion Blur (Platform Shake)	78.45%	0.34%	78.72%	21.21%
Defocus Blur (Out of Focus)	49.16%	3.37%	50.17%	48.15%

BioDINO-Edge is stable under natural illumination changes, keeping **90.24%** accuracy under intense specular sunlight and **91.92%** under harsh canopy shadows. This indicates that ℓ_2 normalization and regularized Fisher covariance shrinkage together decorrelate lighting gradients effectively.

Behavior under optical defocus blur. Under severe defocus blur, raw accuracy falls to 49.16%, while gate rejection stays at 3.37% and the silent error rate reaches 48.15%. Severe defocus physically attenuates the high-frequency structures of fungal lesions (such as pycnidial fruiting specks, concentric necrotic rings, and aecial pustules) into smooth color gradients. The visual backbone can still detect a general foliar canopy, so the softmax posterior may remain overconfident toward visually broader classes, and the Shannon gate does not fire at the default threshold ($\tau_{\text{conf}} = 0.35$). Raising τ_{conf} to 0.60 increases rejection to over 35%, but it brings trade-offs under clean capture. For practical deployment on mobile phones and autonomous UAVs, a lightweight, parameter-free blur detection heuristic, such as a Laplacian variance filter ($\text{Var}(\nabla^2 I) < \theta_{\text{blur}}$, executing in < 0.5 ms), can be placed before neural network inference. It would warn the operator to steady the device or refocus the camera before any compute is spent.

IV. DISCUSSION AND PRACTICAL DEPLOYMENT

Deploying vision AI in real agricultural environments means balancing diagnostic accuracy, hardware latency, and operational autonomy.

Trustworthy robotic scouting. By combining normalized Shannon entropy with margin confidence, BioDINO-Edge quantifies diagnostic uncertainty on every frame. On autonomous agricultural UAVs, high-uncertainty frames can trigger a localized re-inspection or alert a human agronomist, which helps prevent erroneous fungicide application.

The host-context loss dilemma (lesion cropping ablation). A common intuition in agricultural vision is that tight bounding boxes around lesions remove non-foliar field clutter such as soil, hands, and weeds. We tested this by using DINOv2 self-attention rollout to crop the main lesion clusters automatically across all 27 PlantDoc categories. Contrary to expectation, 27-way accuracy dropped slightly from 58.73% to 57.14% (macro F1: 57.95% $\rightarrow$ 56.79%). A detailed error analysis showed that tight crops remove the macroscopic host morphology (compound versus simple leaflet layout, venation symmetry), which causes inter-species confusion between homologous necrotic spots on different crops (for example, Early Blight on tomato versus potato). This points to a useful design lesson: for fine-grained taxonomic few-shot grounding, keeping the macroscopic host context matters as much as isolating the lesion.

Zero-gradient, training-free adaptation. Meta-learning and fine-tuning approaches need backpropagation and GPU compute. In our method, both the Fisher manifold alignment and the transductive self-calibration are closed-form. Edge devices can therefore adapt to newly emerging pathogens within milliseconds, without draining the battery or triggering thermal throttling.

Dual-tier practicality for rural farms. Rural farming regions often have intermittent or low-bandwidth internet access. The dual-tier design keeps the core application compact (a 22.73 MB standalone APK packaging the 15.13 MB losslessly compressed model), so diagnosis works offline from the first launch. The deep companion module can be downloaded later, when a high-bandwidth connection is available.

V. CONCLUSION

We presented **BioDINO-Edge**, an adaptive dual-foundation framework that unites self-supervised visual patch modeling (DINOv2 Small) with organismal taxonomic vision-language supervision (BioCLIP ViT-B). By combining high-dimensional Regularized Fisher-Discriminant Alignment with closed-form Soft-Anchored Transductive Self-Calibration and predictive Shannon uncertainty gating, it sets a new benchmark for few-shot plant pathology: **98.72% peak accuracy**, **98.41% peak F1-score**, and a 30-episode mean of **97.37% $\pm$ 0.73%** on the 11-class PlantVillage 5-shot benchmark, with false-negative disease errors held to **0.604%** and cross-crop confusion to **0.011%**. The system fits in a very small edge footprint (a 22.73 MB standalone APK with a 15.13 MB losslessly compressed fast path, plus an 83.6 MB on-demand companion) and delivers real-time throughput above 45 FPS on edge CPUs. This makes it a trustworthy, field-ready solution for automated crop health scouting.

REFERENCES

[1] B. Richard, A. Qi, and B. D. L. Fitt, "Control of crop diseases through integrated crop management to deliver climate-smart farming systems for low and high input crop production," *Plant Pathology*, vol. 71, no. 1, pp. 187–206, 2022.

[2] A. Khakimov, I. Salakhutdinov, A. Omolikov, and S. Utaganov, "Traditional and current-prospective methods of agricultural plant diseases detection: A review," in *IOP Conference Series: Earth and Environmental Science*, vol. 951, no. 1, 2022, p. 012002.

[3] M. Shoaib, B. Shah, S. EI-Sappagh, A. Ali, A. Ullah, F. Alenezi, T. Gechev, T. Hussain, and F. Ali, "An advanced deep learning models-based plant disease detection: A review of recent research," *Frontiers in Plant Science*, vol. 14, p. 1158933, 2023.

[4] W. Shafik, A. Tufail, A. Namoun, L. De Silva, and R. Apong, "A systematic literature review on plant disease detection: motivations, classification techniques, datasets, challenges, and future trends," *IEEE Access*, vol. 11, pp. 59 174–59 203, 2023.

[5] J. Sun, W. Cao, X. Fu, S. Ochi, and T. Yamanaka, "Few-shot learning for plant disease recognition: A review," *Agronomy Journal*, vol. 116, no. 3, pp. 1204–1216, 2024.

[6] C. Finn, P. Abbeel, and S. Levine, "Model-agnostic meta-learning for fast adaptation of deep networks," in *International Conference on Machine Learning (ICML)*, 2017, pp. 1126–1135.

[7] S. Yang, Q. Feng, J. Zhang, W. Yang, W. Zhou, and W. Yan, "From laboratory to field: cross-domain few-shot learning for crop disease identification in the field," *Frontiers in Plant Science*, vol. 15, p. 1434222, 2024.

[8] Y. Zhao, Z. Zhang, N. Wu, Z. Zhang, and X. Xu, "MAFDE-DN4: Improved few-shot plant disease classification method based on deep nearest neighbor neural network," *Computers and Electronics in Agriculture*, vol. 226, p. 109373, 2024.

[9] J. Mu, Q. Feng, J. Yang, J. Zhang, and S. Yang, "Few-shot disease recognition algorithm based on supervised contrastive learning," *Frontiers in Plant Science*, vol. 15, p. 1341831, 2024.

[10] D. Radočaj, P. Radočaj, I. Plaščak, and M. Jurišić, "Evolution of deep learning approaches in UAV-based crop leaf disease detection: A Web of Science review," *Applied Sciences*, vol. 15, no. 19, p. 10778, 2025.

[11] S. Ahmed, M. B. Hasan, T. Ahmed, and M. H. Kabir, "DExNet: Combining observations of domain adapted critics for leaf disease classification with limited data," *arXiv preprint arXiv:2506.18173*, 2025.

[12] S. Kudo, S. Kagiwada, and H. Iyatomi, "Few-shot metric domain adaptation: Practical learning strategies for an automated plant disease diagnosis," *arXiv preprint arXiv:2412.18859*, 2024.

[13] T. Babu, R. R. Nair, B. Balusamy, W. H. Khoh, and J. Nair, "AgriFewNet—a lightweight RGB few-shot learning model for efficient plant disease classification," *Applied Sciences*, vol. 15, no. 23, p. 12787, 2025.

[14] C. Zhao, T. Xu, Z. Zhang, and X. Geng, "Lightweight cross-domain few-shot plant disease recognition through target-domain statistical calibration," *Sensors*, vol. 26, no. 12, p. 3632, 2026.

[15] H. Zhou, X. Li, Y. Nie, H. Yu, H. Hu, and L. Lin, "VSF: A multimodal few-shot plant disease recognition method in agricultural scenarios," *AgriEngineering*, vol. 8, no. 10, p. 407, 2026.

[16] M. I. Sharif, Y. Zhong, M. Z. Sajid, and F. Marinello, "Real-field-ready and digitally sustainable plant disease recognition via federated multimodal edge learning and few-shot domain adaptation," *Agriculture*, vol. 16, no. 9, p. 918, 2026.

[17] M. Oquab, T. Darcet, T. Moutakanni, H. Vo, M. Szafraniec, V. Khalidov, P. Fernandez, D. Haziza, F. Massa, A. El-Nouby *et al.*, "DINOv2: Learning robust visual features without supervision," *Transactions on Machine Learning Research*, 2024.

[18] S. Stevens, J. Wu, M. J. Thompson, E. G. Campolongo, C. H. Song, D. E. Carlyn, L. Dong, W. M. Dahdul, C. Stewart, T. Berger-Wolf *et al.*, "BioCLIP: A vision foundation model for the tree of life," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024, pp. 19 412–19 424.

[19] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark *et al.*, "Learning transferable visual models from natural language supervision," in *International Conference on Machine Learning (ICML)*, 2021, pp. 8748–8763.

[20] K. Zhou, J. Yang, C. C. Loy, and Z. Liu, "Learning to prompt for vision-language models," *International Journal of Computer Vision*, vol. 130, no. 9, pp. 2337–2348, 2022.

[21] J. Snell, K. Swersky, and R. Zemel, "Prototypical networks for few-shot learning," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 30, 2017.

[22] O. Ledoit and M. Wolf, "A well-conditioned estimator for large-dimensional covariance matrices," *Journal of Multivariate Analysis*, vol. 88, no. 2, pp. 365–411, 2004.

[23] M. Boudiaf, Z. I. Masud, J. Rony, J. Dolz, P. Piantanida, and I. Ben Ayed, "Information maximization for few-shot learning," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 33, 2020, pp. 2445–2457.

[24] D. P. Hughes and M. Salathé, "An open-access repository of images on plant health to enable the development of mobile disease diagnostics," *arXiv preprint arXiv:1511.08060*, 2015.

[25] R. Zhang, W. Zhang, R. Fang, P. Gao, K. Li, J. Dai, Y. Qiao, and H. Li, "Tip-adapter: Training-free CLIP-adapter for better vision-language modeling," in *European Conference on Computer Vision (ECCV)*, 2022, pp. 493–510.

[26] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, "LoRA: Low-rank adaptation of large language models," in *International Conference on Learning Representations (ICLR)*, 2022.

[27] D. Singh, N. Jain, P. Jain, P. Kayal, S. Kumawat, and N. Batra, "PlantDoc: A dataset for visual plant disease detection," in *Proceedings of the 7th ACM IKDD CoDS and 25th COMAD*, 2020, pp. 249–253.

[28] M. Caron, H. Touvron, I. Misra, H. Jégou, J. Mairal, P. Bojanowski, and A. Joulin, "Emerging properties in self-supervised vision transformers," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021, pp. 9650–9660.

[29] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On calibration of modern neural networks," in *International Conference on Machine Learning (ICML)*, 2017, pp. 1321–1330.